

# Multi-Core Programming with OpenMP

Clemens Grelck

University of Amsterdam  
Informatics Institute  
Computer Systems Architecture Group

Amsterdam Computer Science Summer School on  
Parallel Programming  
Amsterdam, July 9–13, 2012

# Outline of the Day

- ▶ Multi-core computer architecture (brief)
- ▶ OpenMP — the Basics
- ▶ Lunch Break
- ▶ OpenMP — the Advanced Story

# Generic Multi-Core Architectures

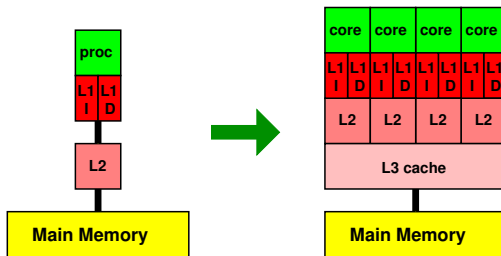
## What we are talking about:

- ▶ Bread and butter multi-core processors
- ▶ Intel XEON family
- ▶ AMD Opteron family
- ▶ maybe Oracle T3/T4 (Niagara)

# Generic Multi-Core Architectures

## Characteristics:

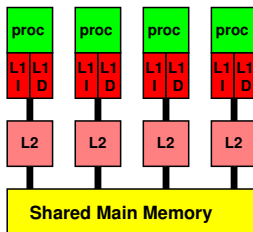
- ▶ Consequence of Murphy's Law: abundance of transistors
- ▶ Put multiple fully-fledged previous generation processors on a single die.
- ▶ Add generous caches
- ▶ Hardware cache coherence



# Generic Multi-Core Architectures

## Symmetric multiprocessing:

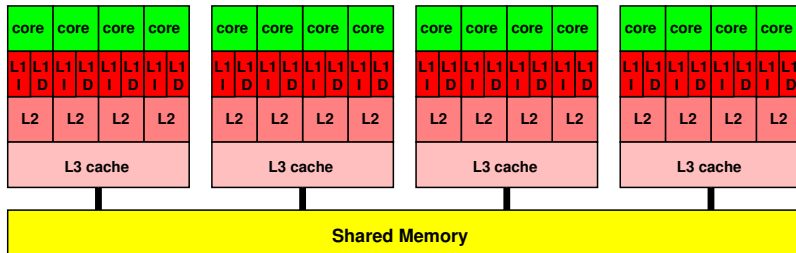
- ▶ Idea dates back to 1980s
- ▶ Assemble multiple processors on one board, 2-way or 4way
- ▶ Processor-local caches
- ▶ Shared main memory
- ▶ Hardware cache coherence



# Generic Multi-Core Architectures

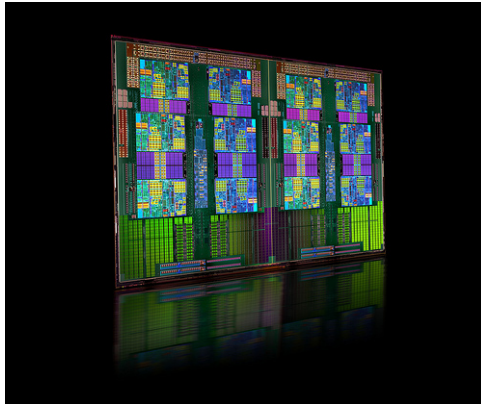
## Multi-core symmetric multiprocessors:

- ▶ Combining symmetric multiprocessing with multi-core processors
- ▶ Unprecedented parallelism on single board
  - ▶ 48 AMD Opteron cores
  - ▶ 24 hyperthreaded Intel Xeon cores
- ▶ Hardware cache coherence



# 12-core AMD Opteron

One chip, two dies:



# A Note on Hyperthreading

## Idea:

- ▶ Duplicate architectural state: (control) registers
- ▶ Do not duplicate functional units
- ▶ Let one core appear as two logical cores to the OS
- ▶ Make more efficient use of resources by computing two streams of instructions simultaneously
- ▶ If one stalls the other takes over



# A Note on Hyperthreading

## Idea:

- ▶ Duplicate architectural state: (control) registers
- ▶ Do not duplicate functional units
- ▶ Let one core appear as two logical cores to the OS
- ▶ Make more efficient use of resources by computing two streams of instructions simultaneously
- ▶ If one stalls the other takes over

## Disadvantage:

- ▶ Requires the right instruction mix
- ▶ Together with symmetric multithreading OS must have specific support for hyperthreading instead of just seeing two cores
- ▶ Cache thrashing

# How to program such systems



# Design Rationale of OpenMP

## Ideal:

- ▶ Automatic parallelisation of sequential code.
- ▶ No additional parallelisation effort for development, debugging, maintenance, etc.

# Design Rationale of OpenMP

## Ideal:

- ▶ Automatic parallelisation of sequential code.
- ▶ No additional parallelisation effort for development, debugging, maintenance, etc.

## Problem:

- ▶ Data dependences are difficult to assess.
- ▶ Compilers must be conservative in their assumptions.

# Design Rationale of OpenMP

## Ideal:

- ▶ Automatic parallelisation of sequential code.
- ▶ No additional parallelisation effort for development, debugging, maintenance, etc.

## Problem:

- ▶ Data dependences are difficult to assess.
- ▶ Compilers must be conservative in their assumptions.

## Way out:

- ▶ Take or write ordinary sequential program.
- ▶ Add annotations/pragmas/compiler directives that guide parallelisation.
- ▶ Let the compiler generate the corresponding code.

# Short History of OpenMP

## Towards OpenMP:

- ▶ Mid 1980s: First commercial SMPs appear.
  - ▶ ANSI standard X3H5 (not adopted)
- ▶ Late 1980s to early 1990s:
  - ▶ Distributed memory architectures prevail.
- ▶ Early 1990s until today:
  - ▶ Shared memory architectures significantly gain popularity.
  - ▶ Different vendor-specific compiler extensions.
- ▶ 1996/97:
  - ▶ Development of a new industry standard: OpenMP
  - ▶ Actively supported by all major hardware vendors: Intel, HP, SGI, IBM, SUN, Compaq, ...
  - ▶ OpenMP Architectural Review Board:  
<http://www.openmp.org/>

# Short History of OpenMP

## OpenMP Standardisation:

- ▶ 1996 : Fortran Version 1.0
- ▶ 1998 : C/C++ Version 1.0
- ▶ 1999 : Fortran Version 1.1
- ▶ 2000 : Fortran Version 2.0
- ▶ 2002 : C/C++ Version 2.0
- ▶ 2005 : Joint Version 2.5
- ▶ 2008 : Joint Version 3.0

# OpenMP at a Glance

## OpenMP as a programming interface:

- ▶ Compiler directives
- ▶ Library functions
- ▶ Environment variables

## C/C++ version:

```
#pragma omp name [clause]*  
structured block
```

## Fortran version:

```
!$ OMP name [clause [, clause]*]  
code block  
!$ OMP END name
```



# Hello World with OpenMP

```
#include "omp.h"
#include <stdio.h>

int main()
{
    printf( "Starting execution with %d threads:\n",
           omp_get_num_threads());

    #pragma omp parallel
    {
        printf( "Hello world says thread %d of %d.\n",
               omp_get_thread_num(),
               omp_get_num_threads());
    }

    printf( "Execution of %d threads terminated.\n",
           omp_get_num_threads());

    return( 0);
}
```

# Hello World with OpenMP

## Compilation:

```
gcc -fopenmp hello_world.c
```

## Output using 4 threads:

```
Starting execution with 1 threads:  
Hello world says thread 2 of 4.  
Hello world says thread 3 of 4.  
Hello world says thread 1 of 4.  
Hello world says thread 0 of 4.  
Execution of 1 threads terminated.
```

# Hello World with OpenMP

## Using 4 threads:

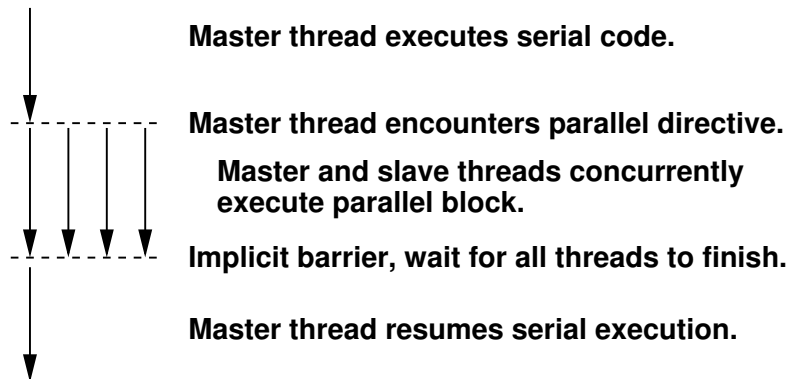
```
Starting execution with 1 threads:  
Hello world says thread 2 of 4.  
Hello world says thread 3 of 4.  
Hello world says thread 1 of 4.  
Hello world says thread 0 of 4.  
Execution of 1 threads terminated.
```

## Who determines number of threads ?

- ▶ Environment variable: `export OMP_NUM_THREADS=4`
- ▶ Library function: `void omp_set_num_threads( int)`

# OpenMP Execution Model

## Classical Fork/Join:



# Simple Loop Parallelisation

## Example: element-wise vector product:

```
void elem_prod( double *c, double *a, double *b, int len)
{
    int i;

    #pragma omp parallel for

    for (i=0; i<len; i++)
    {
        c[i] = a[i] * b[i];
    }
}
```

# Simple Loop Parallelisation

## Example: element-wise vector product:

```
void elem_prod( double *c, double *a, double *b, int len)
{
    int i;

    #pragma omp parallel for

    for (i=0; i<len; i++)
    {
        c[i] = a[i] * b[i];
    }
}
```

## Prerequisite:

- ▶ No data dependence between any two iterations.

# Simple Loop Parallelisation

## Example: element-wise vector product:

```
void elem_prod( double *c, double *a, double *b, int len)
{
    int i;

    #pragma omp parallel for

    for (i=0; i<len; i++)
    {
        c[i] = a[i] * b[i];
    }
}
```

## Prerequisite:

- ▶ No data dependence between any two iterations.
- ▶ **Caution: YOU claim this property !!**

# Directive `#pragma omp parallel for`

## What the compiler directive does for you:

- ▶ It starts additional worker threads depending on `OMP_NUM_THREADS`.
- ▶ It divides the iteration space among all threads.
- ▶ It lets all threads execute loop restricted to their mutually disjoint subsets.
- ▶ It synchronizes all threads at an implicit barrier.
- ▶ It terminates worker threads.



# Directive `#pragma omp parallel for`

## What the compiler directive does for you:

- ▶ It starts additional worker threads depending on `OMP_NUM_THREADS`.
- ▶ It divides the iteration space among all threads.
- ▶ It lets all threads execute loop restricted to their mutually disjoint subsets.
- ▶ It synchronizes all threads at an implicit barrier.
- ▶ It terminates worker threads.

## Restrictions:

- ▶ The directive must directly precede for-loop.
- ▶ The for-loop must match a constrained pattern.
- ▶ The trip-count of the for-loop must be known in advance.

# Shared and Private Variables

## Example:

```
#pragma omp parallel for
for (i=0; i<len; i++)
{
    res[i] = a[i] * b[i];
}
```

- ▶ Shared variable: one instance for **all** threads
- ▶ Private variable: one instance for **each** thread

# Shared and Private Variables

## Example:

```
#pragma omp parallel for  
  
for (i=0; i<len; i++)  
{  
    res[i] = a[i] * b[i];  
}
```

**Who decides that res, a, b, and len are shared variables, whereas i is private ??**

# Shared and Private Variables

## Example:

```
#pragma omp parallel for  
  
for (i=0; i<len; i++)  
{  
    res[i] = a[i] * b[i];  
}
```

Who decides that `res`, `a`, `b`, and `len` are shared variables, whereas `i` is private ??

## Default rules:

- ▶ All variables are **shared**.
- ▶ Only loop variables of parallel loops are **private**.

# Parallelisation of a Less Simple Loop

## Mandelbrot set:

```
double x, y;
int i, j, max = 200;
int depth[M,N];
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    } }
```

# Parallelisation of a Less Simple Loop

## Mandelbrot set:

```
double x, y;
int i, j, max = 200;
int depth[M,N];
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    } }
```

## Properties to check:

- ▶ No data dependencies between loop iterations ?

# Parallelisation of a Less Simple Loop

## Mandelbrot set:

```
double x, y;
int i, j, max = 200;
int depth[M,N];
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    } }
```

## Properties to check:

- ▶ No data dependencies between loop iterations ? **YES !**
- ▶ Trip-count known in advance ?

# Parallelisation of a Less Simple Loop

## Mandelbrot set:

```
double x, y;
int i, j, max = 200;
int depth[M,N];
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    } }
```

## Properties to check:

- ▶ No data dependencies between loop iterations ? **YES !**
- ▶ Trip-count known in advance ? **YES !**
- ▶ Function `mandelval` without side-effects ?



# Parallelisation of a Less Simple Loop

## Mandelbrot set:

```
double x, y;
int i, j, max = 200;
int depth[M,N];
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    } }
```

## Properties to check:

- ▶ No data dependencies between loop iterations ? **YES !**
- ▶ Trip-count known in advance ? **YES !**
- ▶ Function `mandelval` without side-effects ? **YES !**
- ▶ Only loop variable `i` needs to be private ?

# Parallelisation of a Less Simple Loop

## Mandelbrot set:

```
double x, y;
int i, j, max = 200;
int depth[M,N];
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    } }
```

## Properties to check:

- ▶ No data dependencies between loop iterations ? **YES !**
- ▶ Trip-count known in advance ? **YES !**
- ▶ Function `mandelval` without side-effects ? **YES !**
- ▶ Only loop variable `i` needs to be private ? **NO !!!!**

Check `x,y,j`

# Parallelisation of a Less Simple Loop

## Mandelbrot set:

```
double x, y;
int i, j, max = 200;
int depth[M,N];
...
#pragma omp parallel for private( x, y, j)
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    } }
```

## Private clause:

- ▶ Directives may be refined by *clauses*.
- ▶ Private clause allows to tag any variable as private.
- ▶ **Caution:** private variables are **not** initialised outside parallel section !!

# Parallelisation of a Less, Less Simple Loop

## Mandelbrot set with additional counter:

```
int total = 0;
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
        total = total + depth[i,j];
    }
}
```

# Parallelisation of a Less, Less Simple Loop

## Mandelbrot set with additional counter:

```
int total = 0;
...
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
        total = total + depth[i,j];
    }
}
```

## Problems:

- ▶ New variable `total` introduces data dependence.
- ▶ Data dependence could be ignored due to associativity.
- ▶ New variable `total` must be shared.
- ▶ Incrementation of `total` must avoid *race condition*.

# What is a race condition ?

**Note: Adding two values is not atomic**

**Thread 1:**

```
...  
load  reg1 <= &depth  
load  reg2 <= &total  
add    reg1, reg2 => reg1  
store  reg1 => &total  
...
```

# What is a race condition ?

**Note: Adding two values is not atomic**

## Thread 1:

```
...  
load  reg1 <= &depth  
load  reg2 <= &total  
add   reg1, reg2 => reg1  
store reg1 => &total  
...
```

## Thread 2:

```
...  
load  reg1 <= &depth  
load  reg2 <= &total  
add   reg1, reg2 => reg1  
store reg1 => &total  
...
```

## Problem:

- ▶ Threads operate with any interleaving.
- ▶ If different interleavings cause different results, we have a **race condition**.
- ▶ Program behaves non-deterministically.

# What is a race condition ?

**Note: Adding two values is not atomic**

## Thread 1:

```
...  
load  reg1 <= &depth  
load  reg2 <= &total  
add    reg1, reg2 => reg1  
store  reg1 => &total  
...
```

## Thread 2:

```
...  
load  reg1 <= &depth  
load  reg2 <= &total  
add    reg1, reg2 => reg1  
store  reg1 => &total  
...
```

## Solution:

- ▶ We must identify critical regions: sequences of instructions that must be executed without interleaving.
- ▶ We must establish mutual exclusion between critical regions.



# Parallelisation of a Less, Less Simple Loop

## Mandelbrot set with additional counter:

```
int total = 0;
...
#pragma omp parallel for private( x, y, j)
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);

        #pragma omp critical
        {
            total = total + depth[i,j];
        }
    }
}
```

# Critical Regions

## The critical directive:

- ▶ Directive must immediately precede new statement block.
- ▶ Statement block is executed without interleaving.
- ▶ Directive implements critical region.

## Equivalence:

```
#pragma omp critical  
{  
    <statements>  
}
```



```
pthread_mutex_lock( &lock);  
    <statements>  
pthread_mutex_unlock( &lock);
```

# Critical Regions

## The critical directive:

- ▶ Directive must immediately precede new statement block.
- ▶ Statement block is executed without interleaving.
- ▶ Directive implements critical region.

## Equivalence:

```
#pragma omp critical
{
    <statements>
}
```



```
pthread_mutex_lock( &lock);
    <statements>
pthread_mutex_unlock( &lock);
```

## Disadvantage:

- ▶ All critical regions in entire program are synchronised.
- ▶ Unnecessary overhead.

# Critical Regions

## The named critical directive

- ▶ Critical regions may be associated with names.
- ▶ Critical regions with identical names are synchronised.
- ▶ Critical regions with different names are executed concurrently.

## Equivalence:

```
#pragma omp critical (name)
{
    <statements>
}
```



```
pthread_mutex_lock( &name_lock);
    <statements>
pthread_mutex_unlock( &name_lock);
```

# Reduction Operations

## Specific solution: reduction clause

```
#pragma omp parallel for private( x, y, i, j) \  
                                reduction(+:total)  
for (i=0; i<M; i++) {  
    for (j=0; j<N; j++) {  
        x = (double) i / (double) M;  
        y = (double) j / (double) N;  
        depth[i,j] = mandelval( x, y, max);  
        total = total + depth[i,j];  
    }  
}
```

## Properties:

- ▶ Reduction clause only supports built-in reduction operations:  
+, \*, ^, &, |, &&, ||.
- ▶ User-defined reductions only supported via critical regions.
- ▶ Bit accuracy not guaranteed.

# Shared and Private Variables Reloaded

## Shared variables:

- ▶ One instance shared between sequential and parallel execution.
- ▶ Value unaffected by transition.

## Private variables:

- ▶ One instance during sequential execution.
- ▶ One instance per worker thread during parallel execution.
- ▶ No exchange of values.

# Shared and Private Variables Reloaded

## Shared variables:

- ▶ One instance shared between sequential and parallel execution.
- ▶ Value unaffected by transition.

## Private variables:

- ▶ One instance during sequential execution.
- ▶ One instance per worker thread during parallel execution.
- ▶ No exchange of values.

## New: Firstprivate variables:

- ▶ Like private variables, but ...
- ▶ Worker thread instances initialised with master thread value.

# Shared and Private Variables Reloaded

## Shared variables:

- ▶ One instance shared between sequential and parallel execution.
- ▶ Value unaffected by transition.

## Private variables:

- ▶ One instance during sequential execution.
- ▶ One instance per worker thread during parallel execution.
- ▶ No exchange of values.

## New: Firstprivate variables:

- ▶ Like private variables, but ...
- ▶ Worker thread instances initialised with master thread value.

## New: Lastprivate variables:

- ▶ Like private variables, but ...
- ▶ Master thread instance updated to value of worker thread instance that executes the last (in sequential order) iteration.



# Shared and Private Variables Reloaded

## Example:

```
int a=1, b=2, c=3, d=4, e=5;

#pragma omp parallel for private( a) \
                        firstprivate( b,d) \
                        lastprivate( c,d)
for (i=0; i<10; i++) {
    //  a : undef | b : 2 | c : undef | d : 4 | e: 5
    a++; b++; c=7; d++;
}

//  a : 1 | b : 2 | c : 7 | d : ?? | e: 5
```

## Note:

- ▶ The value of d depends on the number of iterations executed by the thread that is assigned iteration i=9 by the OpenMP loop scheduler.

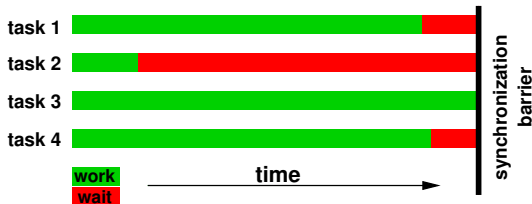
# Loop Scheduling

## Definition:

- ▶ Loop scheduling determines which iterations are executed by which thread.

## Aim:

- ▶ Equal workload distribution



# Loop Scheduling

## Problem:

- ▶ Different situations require different techniques

## The `schedule` clause:

```
#pragma omp parallel for schedule( <type> [, <chunk>])  
for (...)  
{  
    ...  
}
```

## Properties:

- ▶ Clause selects one out of a set of scheduling techniques.
- ▶ Optionally, a chunk size can be specified.
- ▶ Default chunk size depends on scheduling technique.

# Loop Scheduling

## Static scheduling:

```
#pragma omp parallel for schedule( static)
```

- ▶ Loop is subdivided into as many chunks as threads exist.
- ▶ Often called **block scheduling**.

# Loop Scheduling

## Static scheduling:

```
#pragma omp parallel for schedule( static)
```

- ▶ Loop is subdivided into as many chunks as threads exist.
- ▶ Often called **block scheduling**.

## Static scheduling with chunk size:

```
#pragma omp parallel for schedule( static, <n>)
```

- ▶ Loop is subdivided into chunks of  $n$  iterations.
- ▶ Chunks are assigned to threads in a round-robin fashion.
- ▶ Also called **block-cyclic scheduling**.

# Loop Scheduling

## Dynamic scheduling:

```
#pragma omp parallel for schedule( dynamic, <n>)
```

- ▶ Loop is subdivided into chunks of  $n$  iterations.
- ▶ Chunks are dynamically assigned to threads on their demand.
- ▶ Also called **self scheduling**.
- ▶ Default chunk size: 1 iteration.

## Properties:

- ▶ Allows for dynamic load distribution and adjustment.
- ▶ Requires additional synchronization.
- ▶ Generates more overhead than static scheduling.

# Loop Scheduling

## Dilemma of chunk size selection:

- ▶ Small chunk sizes mean good load balancing, but high synchronisation overhead.
- ▶ Large chunk sizes reduce synchronisation overhead, but result in poor load balancing.

# Loop Scheduling

## Dilemma of chunk size selection:

- ▶ Small chunk sizes mean good load balancing, but high synchronisation overhead.
- ▶ Large chunk sizes reduce synchronisation overhead, but result in poor load balancing.

## Rationale of guided scheduling:

- ▶ In the beginning, large chunks keep synchronisation overhead small.
- ▶ When approaching the final barrier, small chunks balance workload.



# Loop Scheduling

## Guided scheduling:

```
#pragma omp parallel for schedule( guided, <n>)
```

- ▶ Chunks are dynamically assigned to threads on their demand.
- ▶ Initial chunk size is implementation dependent.
- ▶ Chunk size decreases exponentially with every assignment.
- ▶ Also called **guided self scheduling**.
- ▶ Minimum chunk size:  $n$  (default: 1)

## Example:

- ▶ Total number of iterations: 250
- ▶ Initial / minimal chunk size: 50 / 5
- ▶ Current chunk size: 80% of last chunk size:  
50 – 40 – 32 – 26 – 21 – 17 – 14 – 12 – 10 – 8 – 6 – 5 – 5 – 4

# Conditional Parallelisation

## Problem:

- ▶ Parallel execution of a loop incurs overhead:
  - ▶ creation of worker threads
  - ▶ scheduling
  - ▶ synchronisation barrier
- ▶ This overhead must be outweighed by sufficient workload.
- ▶ Workload depends on
  - ▶ loop body,
  - ▶ trip count.

# Conditional Parallelisation

## Problem:

- ▶ Parallel execution of a loop incurs overhead:
  - ▶ creation of worker threads
  - ▶ scheduling
  - ▶ synchronisation barrier
- ▶ This overhead must be outweighed by sufficient workload.
- ▶ Workload depends on
  - ▶ loop body,
  - ▶ trip count.

## Example:

```
if (len < 1000) {  
    for (i=0; i<len; i++)  
    {  
        res[i] = a[i] * b[i];  
    }  
}  
else {  
    #pragma omp parallel for  
    for (i=0; i<len; i++)  
    {  
        res[i] = a[i] * b[i];  
    }  
}
```

# Conditional Parallelisation

## Introducing the if-clause:

```
if (len < 1000) {  
    for (i=0; i<len; i++) {  
        res[i] = a[i] * b[i];  
    }  
}  
else {  
    #pragma omp parallel for  
    for (i=0; i<len; i++) {  
        res[i] = a[i] * b[i];  
    }  
}
```



```
#pragma omp parallel for if (len >= 1000)  
for (i=0; i<len; i++) {  
    res[i] = a[i] * b[i];  
}
```

Eventually

Questions



Coffee

Practicals

Lunch