

Multi-Core Programming with OpenMP

Outline:

- Multi-core computer architecture (brief)
- OpenMP — the Basics
- Lunch Break
- OpenMP — the Advanced Story



Exploiting Loop-Level Parallelism with OpenMP

So far: Limited to individual loops:

- No relationship between subsequent parallelizable loops.
- Parallel execution environment repeatedly set up and terminated.
- Frequent barrier synchronizations
- Limited scope for optimization

Wanted: Larger parallel sections

- Containing several parallel loops
- Containing other work sharing constructs
- Barrier synchronization only as necessary
- Overlapping of different parallel activities



Example: Mandelbrot with Dithering

Sequential code:

```
for (i=0; i<M; i++) {  
    for (j=0; j<N; j++) {  
        x = (double) i / (double) M;  
        y = (double) j / (double) N;  
        depth[i,j] = mandelval( x, y, max);  
    }  
}
```

```
for (i=0; i<M; i++) {  
    for (j=1; j<N-1; j++) {  
        dith[i,j] = 0.5 * depth[i,j]  
                    + 0.25 * depth[i,j-1]  
                    + 0.25 * depth[i,j+1];  
    }  
}
```

Outline:

1. Compute Mandelbrot picture
2. Perform dithering step

Mandelbrot with Dithering

Loop-parallelized code:

```
#pragma omp parallel for private(i,j,x,y)
for (i=0; i<M; i++) {
    for (j=0; j<N; j++) {
        x = (double) i / (double) M;
        y = (double) j / (double) N;
        depth[i,j] = mandelval( x, y, max);
    }
}
```

```
#pragma omp parallel for private(i,j)
for (i=0; i<M; i++) {
    for (j=1; j<N-1; j++) {
        dith[i,j] = 0.5 * depth[i,j]
                    + 0.25 * depth[i,j-1]
                    + 0.25 * depth[i,j+1];
    }
}
```

Disadvantage:


- Between the two parallel loops:
 - All threads hit synchronization barrier.
 - Worker threads are terminated.
 - Worker threads are re-created.
- **Avoidable overhead !!**
- Way out: Decoupling of
 - parallelization
 - work sharing

SPMD-Style Parallel Regions

Introducing the parallel-directive:

```
#pragma omp parallel
{
    printf( "Hello world says thread %d.\n", omp_get_thread_num());
}
```

Effect:



SPMD-Style Parallel Regions

Introducing the parallel-directive:

```
#pragma omp parallel
{
    ...
}
```

Meaning:

- The entire code block following the **parallel**-directive is executed by **all** threads concurrently.
- This includes:
 - creation of worker threads,
 - SPMD-style code execution,
 - barrier synchronization,
 - termination of worker threads.

SPMD-Style Parallel Regions

Introducing the parallel-directive:

```
#pragma omp parallel
{
    ...
}
```

Available clauses:

- `private` (*list of variables*)
- `firstprivate` (*list of variables*)
- `shared` (*list of variables*)
- `reduction` (*operator : list of variables*)
- `if` (*logical expression*)

All clauses behave exactly as with `parallel` for directive.

Work Sharing in Parallel Regions

Problem:

- All threads execute identical code within parallel region.
- It is not very useful to let them do exactly the same.
- Work needs to be explicitly shared among threads.

Divide work based on thread number:


```
#pragma omp parallel private( whoami, nthreads)
{
    nthreads = omp_get_num_threads();
    whoami    = omp_get_thread_num();

    do_the_work( whoami, nthreads);
}
```


Work Sharing Constructs in OpenMP

Loop scheduling with the for-directive:

```
int i;  
...  
#pragma omp parallel private(i)  
{  
  
    #pragma omp for  
  
    for (i=0; i<N; i++) {  
        array[i] = ... ;  
    }  
}
```



```
int i;  
...  
  
#pragma omp parallel for  
  
for (i=0; i<N; i++) {  
    array[i] = ... ;  
}
```

- Work sharing / loop scheduling done implicitly by OpenMP
- Directive **parallel for** is equivalent to nested separate directives **parallel** and **for**
- Caution: Loop variables needs explicit privatization

Parallelism vs Work Sharing

The parallel directive:

- defines a parallel region
- starts team of worker threads
- performs a synchronization barrier
- terminates team of worker threads

The for directive:

- does **not** start / stop any threads
- distributes loop iterations among (already running) threads
- **within** a parallel region
- performs a synchronization barrier

Again: Mandelbrot with Dithering

Region-parallelized code:

```
#pragma omp parallel private(i,j,x,y)
{
    #pragma omp for
    for (i=0; i<M; i++) {
        for (j=0; j<N; j++) {
            x = (double) i / (double) M;
            y = (double) j / (double) N;
            depth[i,j] = mandelval( x, y, max);
        }
    }

    #pragma omp for
    for (i=0; i<M; i++) {
        for (j=1; j<N-1; j++) {
            dith[i,j] = 0.5 * depth[i,j]
                + 0.25 * depth[i,j-1]
                + 0.25 * depth[i,j+1];
        }
    }
}
```

Advantages:

- Larger parallel region
- Less thread termination/recreation overhead
- Less communication overhead
- Opportunity for further optimization...

Again: Mandelbrot with Dithering

Region-parallelized code:

```
#pragma omp parallel private(i,j,x,y)
{
    #pragma omp for nowait schedule(static)
    for (i=0; i<M; i++) {
        for (j=0; j<N; j++) {
            x = (double) i / (double) M;
            y = (double) j / (double) N;
            depth[i,j] = mandelval( x, y, max);
        }
    }

    #pragma omp for nowait schedule(static)
    for (i=0; i<M; i++) {
        for (j=1; j<N-1; j++) {
            dith[i,j] = 0.5 * depth[i,j]
                        + 0.25 * depth[i,j-1]
                        + 0.25 * depth[i,j+1];
        }
    }
}
```

New clause: nowait:

- Effect: no synchronization barrier after work sharing construct

Advantages:

- Avoid unnecessary synchronization overhead
- After first parallel loop:
No need to wait for any other thread to finish due to restricted data dependence
- After second parallel loop:
Synchronization done by parallel section anyways

Non-Loop Work Sharing: Parallel Sections

Example:

```
#pragma omp parallel
{
    ...
    #pragma omp sections
    {
        #pragma omp section
        { ...
        }
        #pragma omp section
        { ...
        }
        ...
    }
    ...
}
```

How it works:

- Each parallel section is executed by exactly one thread.
- Threads execute different (maybe unrelated) code
- Mapping of threads to sections is done implicitly.
- $\#threads > \#sections$:
Some threads do nothing.
- $\#threads < \#sections$:
Some threads do multiple sections.
- Synchronization barrier completes **sections** directive.
- The **parallel** directive and the **sections** directive can be combined into the **parallel sections** directive (in analogy to the **parallel for** directive).

Parallel Sections

Clauses of sections directive:

- `private` (*list of variables*)
 - private copy
 - not initialized from surrounding context
- `firstprivate` (*list of variables*)
 - private copy
 - initialized with pre-section value
- `lastprivate` (*list of variables*)
 - private copy
 - value in last section propagated outside
- `reduction` (*operator : list of variables*)
 - private copy
 - each section's value folded with reduction operator and result propagated outside

Parallel Sections Example

Simulation Program:

- Read input data from file.
- Perform three independent preprocessing steps:
 1. Interpolation of input data to regular grid.
 2. Gathering statistics about input data.
 3. Generation of random parameters for Monte Carlo simulation.
- Perform simulation.

```
data = read_input( );
#pragma omp parallel sections
{
    #pragma omp section
    {
        grid = interpolate( data);
    }
    #pragma omp section
    {
        stats = compute_statistics( data);
    }
    #pragma omp section
    {
        rands = generate_rand_params( );
    }
}
simulate( grid);
```

Assigning Work to a Single Thread

The single directive:

- Work sharing construct
- Exactly one thread executes code block
- Synchronization barrier afterwards
- Useful to “interrupt” parallel region
- **nowait** clause leaves out synchronization
- Semantically equivalent but usually more efficient than two separate parallel regions

```
#pragma omp parallel \  
    shared( in, out, len, array)  
{  
    ...  
    #pragma omp single  
    {  
        read_array( in, array, len);  
    }  
    #pragma omp for  
    for (i=0; i<len; i++) {  
        array[i] = fun( array[i]);  
    }  
    #pragma omp single nowait  
    {  
        write_array( out, array, len);  
    }  
    ...  
}
```


Work Sharing Constructs in OpenMP

Restrictions:

- Single point of entry.
 - No jumps into work sharing construct.
- Single point of exit.
 - No jumps out of work sharing construct.
 - No **return** or **break** statements.
- All threads must encounter work sharing construct in same order.
 - No subsets of thread teams.
- No nesting
 - Control reaches work sharing construct while already executing another one:
 - * Programming error.
 - * Behaviour undefined.

Orphaned Work Sharing Constructs

What is that ?

- Work sharing construct outside lexical scope of surrounding parallel region.
- Perfectly fine with OpenMP.
- Work sharing construct is associated with innermost parallel region in dynamic function call tree.

No parallel region around ?

- If work sharing construct is encountered during serial execution, it behaves like when executed by a team of a single thread.

Example:

```
void work( )
{ ...
    #pragma omp parallel
    {
        initialize( array, N);
    } ...
}
```

```
void initialize( int *a, int N)
{
    int i;
    #pragma omp for
    for (i=0; i<N; i++) {
        array[i] = ... ;
    }
}
```

Thread Private Global Variables

Idea:

- Global variables are always shared between all threads.
- Sometimes, it is more convenient to have individual copy for each thread.
- Analogy to *thread specific data (TSD)* in PThreads.
- Thread private instances of global variables survive across different parallel regions.
- Conceptually, there are $N+1$ instances, one for the master thread and one for each worker thread.

Example:

```
#pragma omp threadprivate
int my_id;

int main()
{
    #pragma omp parallel
    {
        my_id = omp_get_thread_num();
        ...
    }
    ...
}
```

Thread Private Global Variables

The copyin clause:

- Usually, there is no way to exchange data stored in thread private variables between threads.
- The **copyin** clause allows to initialize all copies with master thread's value when entering parallel region.

Example:

```
#pragma omp threadprivate
int N, my_id;

int main()
{
    N = ... ;

    #pragma omp parallel copyin (N)
    {
        my_id = omp_get_thread_num();
        ...
        N = fun( N, my_id);
    }
    ...
}
```

Synchronization Barrier in OpenMP

The barrier directive:

- All threads wait at barrier until last thread reaches barrier.
- Allows for explicit control over synchronization barriers independent of work sharing constructs.

Example:

```
#pragma omp parallel
{
    ...
    #pragma omp barrier
    ...
}
```

Low-level Synchronization

Similar API as PThreads:

```
void omp_init_lock( omp_lock_t *lock);  
void omp_destroy_lock(omp_lock_t *lock);  
void omp_set_lock(omp_lock_t *lock);  
void omp_unset_lock(omp_lock_t *lock);  
int omp_test_lock(omp_lock_t *lock);
```

Nested critical sections:

- Same problems as with mutex locks in PThreads.
- May deadlock easily.
- OpenMP provides no specific means for deadlock prevention or detection.



Conclusion

Advantages:

- Simpler to use than other concepts, for example PThreads.
- OpenMP implementation responsible for most organizational aspects.
- Simplicity facilitates experimentation with alternatives, e.g. scheduling techniques.
- Large applications may be parallelized incrementally.

Disadvantages:

- Simplicity reduces programmers' reflections on parallelism.
- Insight into impact of certain directives / clauses is reduced.
- False directives may lead to wrong — even non-deterministic — behaviour.
- False directives are easy to write but hard to find.
- Performance considerations lead to low-level “expert” code.

Outlook: OpenMP 3.0

Extensions:

- Nested parallel regions:
 - Parallel regions may be nested lexically or dynamically
 - Additional level of parallelism
 - New team of threads started
 - Creator thread becomes master of that new team
 - Controlled by **num_threads** clause
- Independent tasks:
 - Spawn concurrent evaluation of some expression
 - Wait on termination of spawned tasks
 - Clauses control visibility of variables
- ...



OpenMP Research

Research projects:

- OpenMP for the Cell
- OpenMP for GPGPUs
- Extension for OpenMP for GPGPUs
- ...



The End

Questions ?

More information:

www.openmp.org

Multi-Core Programming with OpenMP



Dr. Clemens Grelck
Computer Systems Architecture
University of Amsterdam