

PARALLEL PROGRAMMING

DAY 5
DATAFLOW PROGRAMMING

OVERVIEW

- ✻ Turing machines / imperative programming:

- ✻ “do this, **then** do that” : **time** ordered

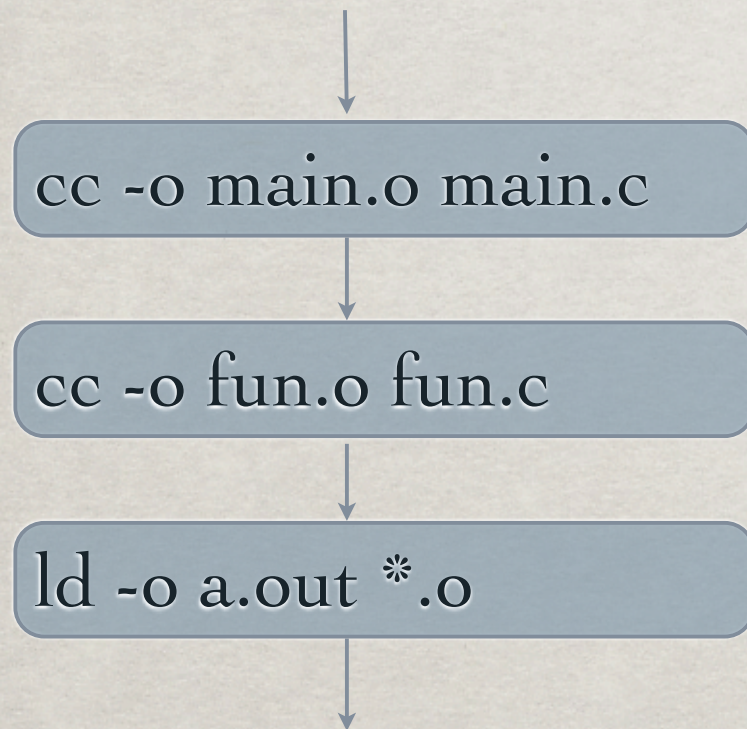
- ✻ instructions “active,” data is “passive”

- ✻ Dataflow programming:


- ✻ “to compute this, first compute that” :
dependency ordered

- ✻ data is “active,” instructions “passive”

EXAMPLE



Imperative program



Dataflow program

FUNCTIONS VS STREAMS

- ✻ Functional and dataflow: both represent the **functional dependencies between values**: almost **no hidden state**
- ✻ What's the difference?
 - ✻ Functional “programs” compute 1 value
(Any more requires read-eval-print or “infinite recursion”)
 - ✻ Dataflow programs (usually) operate over **streams**: “flow” of inputs

WHY IT MATTERS

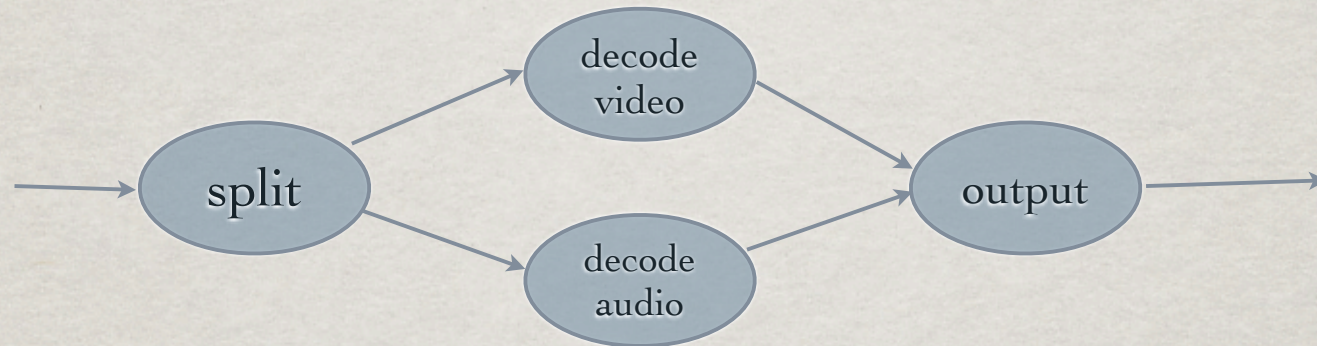
- ✿ **Programmer sanity:** no hidden state means sub-programs can be analyzed/tested individually
- ✿ Relevant to this summer school:
 - ✿ Two operations that have their input ready can run simultaneously: **parallelization is easy**
 - ✿ synchronization needed at “merge” nodes
 - ✿ Inherently **heterogeneous** concurrency

WHERE AND HOW

- ✿ Dataflow is a **style of programming** which can be used in most parallel programming languages:
 - ✿ **identify dependencies** between data and operations
 - ✿ let the **environment schedule execution**
- ✿ Few pure “dataflow programming languages,” most common:
 - ✿ VHDL/Verilog
 - ✿ Spreadsheet formulas


TECHNIQUES

FLAVORS OF DATAFLOW PROGRAMS



- ✿ Synchronous: global time step, all nodes activated on their “current input”
- ✿ Asynchronous: each node computes at its own speed, may suspend on missing input or full output buffer

FLAVORS OF DATAFLOW PROGRAMS



- ✱ Static: the number of nodes and edges stays the same over time
- ✱ Dynamic: nodes and edges appear and disappear over time

EXAMPLE USES

EXAMPLE USES

- ✱ Static, asynchronous:
video compression/decompression

EXAMPLE USES

- ✻ Static, asynchronous:
video compression/decompression
- ✻ Static, synchronous:
video filtering

EXAMPLE USES

- ✱ Static, asynchronous:
video compression/decompression
- ✱ Static, synchronous:
video filtering
- ✱ Dynamic, asynchronous:
any functional computations with recursion

EXAMPLE USES

- ✻ Static, asynchronous:
video compression/decompression
- ✻ Static, synchronous:
video filtering
- ✻ Dynamic, asynchronous:
any functional computations with recursion
- ✻ Dynamic, synchronous: N/A

HOW TO EXECUTE STATIC SYNCHRONOUS PROGRAMS

- ✱ Program = graph
 - ✱ Operations as nodes, data as edges
- ✱ Implement each node as a **function + ready bit**
- ✱ Use a **buffer for each data edge**
- ✱ At each cycle:
 - ✱ execute ready functions
 - ✱ store result into output buffer
 - ✱ mark functions on the other side of buffers as ready

HOW TO EXECUTE ASYNCHRONOUS PROGRAMS

- ✱ Program = graph
 - ✱ Operations as nodes, data as edges
- ✱ At any point during execution, nodes are **activated** based on data availability
- ✱ The environment can **fire** operations in any order, **possibly in parallel**
- ✱ Implementations: **pipeline, job pooling**
- ✱ At low level: **FIFOs and/or dataflow variables**

QUEUES / FIFOs

- ✱ Each operation runs in a process connected to other processes using FIFOs representing data edges
- ✱ “First in, first out”
 - ✱ can be implemented in software as array + pointers to “head” and “tail”
 - ✱ Can be implemented in hardware too
- ✱ Processes **suspend** when:
 - ✱ read from empty: tail reaches head
(resume when data becomes available)
 - ✱ write to full: head reaches tail
(resume when space available in FIFO)

DATAFLOW VARIABLES

- ✿ Each operation is expressed as a function which saves its return value into a dataflow variable
- ✿ **Dataflow variable = data + state**
 - ✿ Data: a value, like usual
 - ✿ **State: full/empty**
- ✿ Initially empty
- ✿ “Read from empty” suspends the current processes and writes the process ID to the dataflow variable
- ✿ Write to suspended wakes up the process(es) whose ID is saved in the variable

HYBRID DATAFLOW

- ✿ Often:
pipes and dataflow variables are **narrow**
 - ✿ Can only pass **a few bytes at a time efficiently** through the device
- ✿ Combine with a **shared data store**
 - ✿ Shared memory for multiple threads
 - ✿ Shared filesystem for Unix processes
- ✿ **Pass the names (references)** to data items through the dataflow device,
- ✿ **Use the shared store for the body** of data items

HIGH-LEVEL REASONING ABOUT DATAFLOW BEHAVIOR

LATENCY

- ⌘ Latency for sub-graphs of a dataflow program:
time between **arrival of input** item and **production of corresponding output** item
- ⌘ Latency constrained by:
 - ⌘ **latency of data links** (secs)
 - ⌘ **performance of computing resource** (secs/op)
 - ⌘ **latency to shared data store** (secs) for hybrid dataflow
- ⌘ General constraint:
 $L > \text{max-flow}(p \odot g)$
 p = program graph (op nodes + dependency edges)
 g = resource graph (PUs + data links)
 $p \odot g$: mapping of program onto resources
max-flow : **sum of latencies along the longest path** through the resource graph
that must be followed to compute 1 output
also called **critical path**

LATENCY OPTIMIZATION

LATENCY OPTIMIZATION

- ✱ Some problems are **latency-constrained**:

LATENCY OPTIMIZATION

- ✱ Some problems are **latency-constrained**:

- ✱ Medical equipment: eg $L < 40\text{ms}$

LATENCY OPTIMIZATION

- ✱ Some problems are **latency-constrained**:
 - ✱ Medical equipment: eg $L < 40\text{ms}$
 - ✱ Weather forecast: eg $L < 1 \text{ day}$

LATENCY OPTIMIZATION

- ✱ Some problems are **latency-constrained**:

- ✱ Medical equipment: eg $L < 40\text{ms}$

- ✱ Weather forecast: eg $L < 1 \text{ day}$

- ✱ “Levers” to reduce latency:

LATENCY OPTIMIZATION

- ✱ Some problems are **latency-constrained**:

- ✱ Medical equipment: eg $L < 40\text{ms}$

- ✱ Weather forecast: eg $L < 1 \text{ day}$

- ✱ “Levers” to reduce latency:

- ✱ faster PUs, faster data links

LATENCY OPTIMIZATION

- ✱ Some problems are **latency-constrained**:

- ✱ Medical equipment: eg $L < 40\text{ms}$

- ✱ Weather forecast: eg $L < 1 \text{ day}$

- ✱ “Levers” to reduce latency:

- ✱ faster PUs, faster data links

- ✱ shorter critical path for the computation

THROUGHPUT

- ✿ Throughput for sub-graphs of a dataflow program:
how many nodes are fired by second
- ✿ Throughput constrained by:
 - ✿ **bandwidth of data links** (bytes/sec)
 - ✿ **max throughput of computing resource** (ops/sec)
 - ✿ **bandwidth of shared data store** (bytes/sec) for hybrid dataflow
- ✿ General constraint:
$$T < \min(\text{Bin}/a, \text{MaxInternalT}, \text{Bout}/b)$$
 - a = number of input bytes consumed per op
 - b = number of output bytes produced per op
 - Bin/Bout = bandwidth of data links where computation is mapped
 - MaxInternalT = maximum internal throughput for the sub-parts

THROUGHPUT OPTIMIZATION

THROUGHPUT OPTIMIZATION

- ✱ Most computing goals are about **throughput maximization**

THROUGHPUT OPTIMIZATION

- ✱ Most computing goals are about **throughput maximization**
- ✱ Video / games: frames / sec

THROUGHPUT OPTIMIZATION

- ✱ Most computing goals are about **throughput maximization**
- ✱ Video / games: frames / sec
- ✱ Build system: compilations / sec

THROUGHPUT OPTIMIZATION

- ✱ Most computing goals are about **throughput maximization**
- ✱ Video / games: frames / sec
- ✱ Build system: compilations / sec
- ✱ “Levers” to increase throughput:

THROUGHPUT OPTIMIZATION

- ✱ Most computing goals are about **throughput maximization**
- ✱ Video / games: frames / sec
- ✱ Build system: compilations / sec
- ✱ “Levers” to increase throughput:
 - ✱ bandwidth of data links

THROUGHPUT OPTIMIZATION

- ✱ Most computing goals are about **throughput maximization**
 - ✱ Video / games: frames / sec
 - ✱ Build system: compilations / sec
- ✱ “Levers” to increase throughput:
 - ✱ bandwidth of data links
 - ✱ Number of parallel PUs

DATAFLOW WITH UNIX PIPES

UNIX PIPES

- ✻ Basic operations of processes: **read/write**
- ✻ Basic interface to the environment: **file names**
- ✻ A kind of special files: FIFOs
 - ✻ effect: direct data link between 2 processes
 - ✻ create: **mkfifo** (named), **pipe** (unnamed)
 - ✻ use: **read/write** - read empty blocks, write resumes suspended process

IMAGE FILTERS++

- ✻ Basic idea: hybrid, static asynchronous
 - ✻ Use **pipes** to pass the **image names**
 - ✻ Use the **filesystem** to pass the **image data**
 - ✻ A name goes through the pipe only after the image data has been committed to file
- ✻ Merge / synchronize:
successive reads to multiple input pipes